

CS 577 - Greedy

Marc Renault

Department of Computer Sciences
University of Wisconsin – Madison

Summer 2023

TopHat Section 001 Join Code: 275653



GREEDY

GREEDY ALGORITHMS

What is a Greedy Algorithm (GREEDY)?

- Typically, thought of as a *heuristic* that is locally optimal.
- Is GREEDY always the best? No, but a good place to start.
- This notion has yet to be fully formalized, and it often problem specific.

Definition from Priority Algorithms

A greedy algorithm is an algorithm that processes the input in a specified order. For each request in the input, the greedy algorithm processes it so as to minimize (resp. maximize) the objective, assuming that the request is the last request.

For a given problem, there may be many greedy algorithms.

IS GREEDY OPTIMAL?

Not always: Bin Packing Problem

- Bins of size 1, and requests of size $(0, 1]$.
- Objective: Pack the items in the minimum number of bins.
- Greedy heuristic: FIRST FIT INCREASING (FFI)

Non-optimal example:


- $\sigma = \langle 1/2 - \varepsilon, 1/2 - \varepsilon, 1/2 + \varepsilon, 1/2 + \varepsilon \rangle$
- FFI: 3 bins
- OPT: 2 bins

Techniques for showing that GREEDY is optimal:

- Always stays ahead
- Exchange argument

STAYS AHEAD: INTERVAL SCHEDULING

INTERVAL SCHEDULING



Problem Definition

- Requests: $\sigma = \{r_1, \dots, r_n\}$
- A request $r_i = (s_i, f_i)$, where s_i is the start time and f_i is the finish time.
- Objective: Produce a *compatible* schedule S that has maximum cardinality.
- Compatible schedule S : $\forall r_i, r_j \in S, f_i \leq s_j \vee f_j \leq s_i$.

TopHat Discussion 1: What greedy heuristic might work?


GREEDY ALGORITHMS FOR INTERVAL SCHEDULING

Heuristic 1: Earliest First

Schedule a compatible request with the earliest start time.

Optimal?

Counter-example:




GREEDY ALGORITHMS FOR INTERVAL SCHEDULING

Heuristic 2: Smallest Interval

Schedule a compatible request r_i with the smallest interval ($f_i - s_i$).

Optimal?

Counter-example:



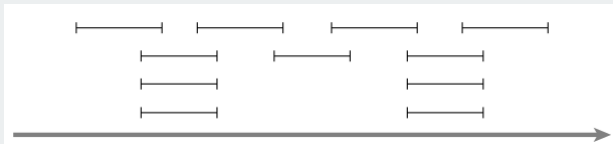
GREEDY ALGORITHMS FOR INTERVAL SCHEDULING

Heuristic 3: Fewest Conflicts

Schedule a compatible request with the fewest remaining conflicts.

Optimal?

Counter-example:



GREEDY ALGORITHMS FOR INTERVAL SCHEDULING

Heuristic 4: Finish First

Schedule a compatible request with the smallest finish time.

Optimal?

Counter-example? Let's try and prove it.

EXERCISE: FORMALIZE THE ALGORITHM (PSEUDOCODE)

HEURISTIC 4: FINISH FIRST

Algorithm: FINISHFIRST

Let S be an initially empty set.

while σ is not empty **do**

 Choose $r_i \in \sigma$ with the smallest finish time (break ties arbitrarily).


 Add r_i to S .

 Remove all incompatible request in σ .

end

return S

Sample Run (TopHat Q1: What is $|S|$?)



ANALYSIS OF FINISHFIRST

Observation 1

Immediate from the definition of FINISHFIRST, S is compatible.

Showing Optimality

Let S^* be an optimal solution.

- We can show the strong claim that $S = S^*$.
- Can there be multiple S^* ? Yes.
- Hence, we can show the weaker claim of $|S| = |S^*|$ for this problem.
- Technique: “Always stays ahead”
 - At every time step i , $|S_i| \geq |S_i^*|$.

STAYS AHEAD ANALYSIS

- Label $S = \langle i_1, \dots, i_k \rangle$ such that $f_{i_u} < f_{i_v}$ for $u < v$.
- Label $S^* = \langle j_1, \dots, j_m \rangle$ such that $f_{j_u} < f_{j_v}$ for $u < v$.

Lemma 1

For all i_r, j_r with $r \leq k$, we have $f_{i_r} \leq f_{j_r}$

Proof.

The proof is by induction.

- For $r = 1$, the claim is true as FINISHFIRST first selects the request with the earliest finish time.
- Assume true for $r - 1$.
 - By the induction hypothesis, we have that $f_{i_{r-1}} \leq f_{j_{r-1}}$.
 - The only way for S to fall behind S^* would be for FINISHFIRST to choose a request q with $f_q > f_{j_r}$, but this is a contradiction.



STAYS AHEAD ANALYSIS

- Label $S = \langle i_1, \dots, i_k \rangle$ such that $f_{i_u} < f_{i_v}$ for $u < v$.
- Label $S^* = \langle j_1, \dots, j_m \rangle$ such that $f_{j_u} < f_{j_v}$ for $u < v$.

Lemma 1

For all i_r, j_r with $r \leq k$, we have $f_{i_r} \leq f_{j_r}$

The optimality of FINISHFIRST, essentially, follows immediately from Lemma 1.

FINISHFIRST IS OPTIMAL

- Label $S = \langle i_1, \dots, i_k \rangle$ such that $f_{i_u} < f_{i_v}$ for $u < v$.
- Label $S^* = \langle j_1, \dots, j_m \rangle$ such that $f_{j_u} < f_{j_v}$ for $u < v$.

Theorem 2

FINISHFIRST produces an optimal schedule.

Proof.

By way of contradiction, assume that $|S^*| > |S|$. This implies that $m > k$. Lemma 1 shows that FINISHFIRST is ahead for all the k requests. That means it would be able to add the $(k + 1)$ -st item of S^* . As it did not, this contradicts the definition of FINISHFIRST. □

IMPLEMENTATION AND RUNNING TIME

Algorithm: FINISHFIRST

Let S be an initially empty set.

while σ is not empty **do**

 Choose $r_i \in \sigma$ with the smallest finish time (break ties arbitrarily).

 Add r_i to S .

 Remove all incompatible request in σ .

end

return S

Implementation Details

- Choose request with smallest finish time: Before processing, sort requests: $O(n \log n)$.
- Remove incompatible requests: Advance in sorted order until a request with a compatible start time.

Overall:

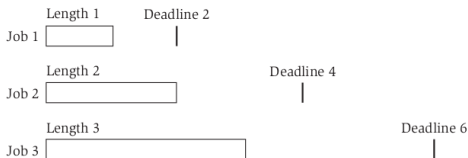
$$O(n \log n) + O(n) = O(n \log n)$$

INTERVAL EXTENSIONS

- Online variant: Requests are presented in a specific order to the algorithm. At request i , the algorithm does not know n nor r_{i+1}, \dots, r_n .
- Add a value to the intervals (online/offline). Now objective is to maximize the total value of scheduled intervals.
- Scheduling all intervals: Interval Colouring Problem.
 - Unlimited resources and the algorithm must produce multiple compatible schedules that cover all the requests (without duplicates between the schedules).
 - Objective: Minimize the number of schedules.

EXCHANGE ARGUMENT: MINIMIZE MAX LATENESS

SCHEDULING PROBLEM: MINIMIZE LATENESS



Problem Definition

- n jobs and a single machine that can process one job at a time
- For job i :
 - t_i is the processing time, d_i is the deadline.
 - Lateness $l_i = f_i - d_i$ if finish time $f_i > d_i$; 0 otherwise.
- Objective: Build a schedule for all the jobs that minimizes the max lateness.

TopHat Discussion 2: What greedy heuristic might work?

GREEDY ALGORITHMS FOR MINIMIZING MAX LATENESS

Heuristic 1: Increasing processing time.

Schedule jobs by increasing t_i .

Optimal?

Counter-example: Jobs (t_i, d_i) : $\{(1, 100), (10, 10)\}$

GREEDY ALGORITHMS FOR MINIMIZING MAX LATENESS

Heuristic 2: Increasing slack.

Schedule by increasing $d_i - t_i$.

Optimal?

Counter-example:

Jobs (t_i, d_i) : $\{(1, 2), (10, 10)\}$

GREEDY ALGORITHMS FOR MINIMIZING MAX LATENESS

Heuristic 3: Earliest deadline first.

Schedule by increasing d_i .

Optimal?

Counter-example? Let's try and prove it.

EXERCISE: FORMALIZE THE ALGORITHM (PSEUDOCODE)

HEURISTIC 3: EARLIEST DEADLINE FIRST.

Algorithm: EDF

Let J be the set of jobs.

Let S be an initially empty list.

while J is not empty **do**


 Choose $j \in J$ with the smallest d_i (break ties arbitrarily).

 Append j to S .

end

return S

Sample Run (TopHat Q1: What is max lateness?)



ANALYSIS OF EDF

Observation 2

There is an optimal schedule with no idle time.

Showing Optimality

Let S^* be an optimal solution.

- Is it sufficient to show that $|S| = |S^*|$? No.
- Can there be multiple S^* ? Yes.
- We need to show either $S = S^*$, or $S \equiv S^*$ for max lateness.
- Technique: “Exchange Argument”
 - Start with an optimal solution S^* and transform it over a series of steps to something equivalent to S while maintaining optimality.
 - $S^* \equiv S_1 \equiv S_2 \equiv \dots \equiv S$ for max lateness.

EXCHANGE ARGUMENT ANALYSIS

Definition 3

A schedule A has an *inversion* if there are jobs i and j with i scheduled before j and $d_j < d_i$.

Lemma 4

All schedules with no inversions and no idle time have the same lateness.

Proof.

- Only vary in jobs with the same deadline.
- Jobs with same deadline must be sequential.
- Ordering of jobs with same deadline won't change lateness.



ANALYSIS OF EDF

Theorem 5

There is an optimal schedule that has no inversions and no idle time.

Proof.

- If S^* has an inversion, then there is a pair of jobs i and j such that j is scheduled immediately after i and has $d_j < d_i$.
- We will swap i and j to create a new schedule S' . Note that S' has one less inversion than S^* .
- We need to show that S' has the same max lateness as S^* :
 - Swapping i and j means that l'_j (lateness in S') is less than that in S^* .
 - Lateness of i may increase, but:
$$l'_i = f'_i - d_i = f_j^* - d_i \leq f_j^* - d_j = l_j^*.$$
- Let $S^* := S'$ and repeat until no more inversions.



EDF IS OPTIMAL

Corollary 6

EDF produces an optimal schedule.

Proof.

- EDF produces a schedule with no inversions and no idle time.
- From Theorem 5, there is an optimal schedule with no inversions and no idle time.
- Lemma 4 shows that these two schedules have the same max lateness.



Run time: Sort the jobs by deadline: $O(n \log n)$.


SHORTEST PATH

FINDING THE SHORTEST PATH

Problem Definition

We have a directed graph $G = (V, E)$, where $|V| = n$ and $|E| = m$ and a node s that has a path to every other node in V . For each edge e , $\ell_e \geq 0$ is the length of the edge.

- What is the shortest path from s to each other node?



Edsger Dijkstra, 1956
Dijkstra's shortest path fame

DIJKSTRA'S

Algorithm: *Dijkstra's*

Let S be the set of explored nodes.

For each $u \in S$, we store a distance value $d(u)$.

Initialize $S = \{s\}$ and $d(s) = 0$

while $S \neq V$ **do**

 Choose $v \notin S$ with at least one incoming edge
 originating from a node in S with the smallest

$$d'(v) = \min_{e=(u,v):u \in S} \{d(u) + \ell_e\}$$

 Append v to S and define $d(v) = d'(v)$.

end

How is it greedy?

TopHat 3: Which technique to prove optimality?

CORRECTNESS OF DIJKSTRA'S

Theorem 7

Consider the S at any point in the execution of Dijkstra's. For each $u \in S$, the path P_u is a shortest $s - u$ path.

Proof.

By induction on the size of S .

- For $|S| = 1$, the claim follows trivially as $S = \{s\}$.
- By the induction hypothesis, for $|S| = k$, P_u is the shortest $s - u$ path for all $u \in S$.

CORRECTNESS OF DIJKSTRA'S

Theorem 7

Consider the S at any point in the execution of Dijkstra's. For each $u \in S$, the path P_u is a shortest $s - u$ path.

Proof.

By induction on the size of S .

- In step $k + 1$, we add v .
 - By definition, P_v is shortest path connected to S by one edge.
 - Since P_u is a shortest path to u , P_v is the shortest path to v when considering only the nodes of S .
 - Moreover, there cannot be a shorter path to v passing through another node $y \notin S$ else y that would be added at $k + 1$.



DIJKSTRA'S OBSERVATIONS

Algorithm: *Dijkstra's*

Let S be the set of explored nodes.
For each $u \in S$, we store a distance
value $d(u)$.

Initialize $S = \{s\}$ and $d(s) = 0$

while $S \neq V$ **do**

 Choose $v \notin S$ with at least one
 incoming edge originating
 from a node in S with the
 smallest $d'(v) =$

$\min_{e=(u,v):u \in S} \{d(u) + \ell_e\}$

 Append v to S and define
 $d(v) = d'(v)$.

end

- Negative edge weights, where does it fail?
- TopHat 4: It is graph exploration, what kind of exploration?
 - Weighted (continuous) BFS

IMPLEMENTATION AND RUN TIME OF DIJKSTRA'S

Algorithm: *Dijkstra's*

Let S be the set of explored nodes.
For each $u \in S$, we store a distance
value $d(u)$.

Initialize $S = \{s\}$ and $d(s) = 0$

while $S \neq V$ **do**

 Choose $v \notin S$ with at least one
 incoming edge originating
 from a node in S with the
 smallest $d'(v) =$

$\min_{e=(u,v):u \in S} \{d(u) + \ell_e\}$



 Append v to S and define
 $d(v) = d'(v)$.

end

- TopHat 5:
Number of
iterations of
the loop?
 $n - 1$
- Key
Operations:
 - Finding
the min:
Easy in
 $O(m)$
- Overall:
 $O(mn)$
- How can we
get
 $O(m \log n)$?

PAGING

PAGING PROBLEM



Cache:



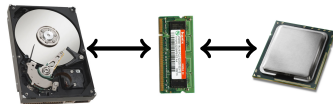

Requests:



Definition

- \mathcal{U} : universe of pages ($|\mathcal{U}| > k$).
- Cache of size k .
- Requests are to the pages of \mathcal{U} .
- Goal: Minimize the number of page faults (requests to pages not in the cache).

PAGING PROBLEM



Cache:



Requests:



Eviction Strategies

- When designing an algorithm, we are picking an eviction strategy.
- In the offline version, the algorithm knows the request sequence. What might be a good eviction strategy?

OFFLINE GREEDY ALGORITHM

Farthest-in-Future (FF)

Evict the page whose next request is the furthest into the future.

Small Run:

- $\mathcal{U} = \{a, b, c\}$
- $k = 2$
- $\sigma = \langle a, b, c, b, c, a, b \rangle$
- TopHat 6: How many faults in small run?

TopHat 7: Which strategy to prove optimality?

PROVING FF OPTIMALITY

EXCHANGE ARGUMENT

Theorem 8

Let S be a schedule for the n request that make the same eviction decisions as S_{FF} for the first j items. Then, there is a schedule S' that makes the same eviction requests as S_{FF} for the first $j + 1$ items with no more faults than S .

Proof.

- If on request $j + 1$, S behaves as S_{FF} . Then define S' as S and the claim follows.
- Otherwise, say S evicts u and S_{FF} evicts v . We will build S' by following S_{FF} for the first $j + 1$ requests. Note that the number of faults are the same for S and S' up to $j + 1$, and the caches match except for u and v .

PROVING FF OPTIMALITY

EXCHANGE ARGUMENT

Theorem 8

Let S be a schedule for the n request that make the same eviction decisions as S_{FF} for the first j items. Then, there is a schedule S' that makes the same eviction requests as S_{FF} for the first $j + 1$ items with no more faults than S .

Proof.

- From $j + 2$ onward, S' follows S until either:
 - ① S evicts v . In this case, S' evicts u .
 - ② S evicts $g \neq v$ to bring u into the cache. In this case, S' evicts g and brings in v .
 - Note: Since S_{FF} evicts v at $j + 1$, u must be requested before v .
- In either case, both S and S' have a page fault, and afterwards their cache match.



PROVING FF OPTIMALITY

EXCHANGE ARGUMENT

Theorem 8

Let S be a schedule for the n request that make the same eviction decisions as S_{FF} for the first j items. Then, there is a schedule S' that makes the same eviction requests as S_{FF} for the first $j + 1$ items with no more faults than S .

How do we get optimality of S_{FF} from Theorem 8?

By induction: We begin with the optimal schedule S^* and inductively apply Theorem 8 for $j = 1, 2, 3, \dots, n$, which after the n iterations, produces S_{FF} .

MST

MINIMUM SPANNING TREE PROBLEM

MST Problem

Let $G = (V, E)$ be a connected graph, where $|V| = n$ and $|E| = m$. For each edge e , $c_e > 0$ is the cost of the edge.

- Find an edge set $F \subseteq E$ with minimum cost that keeps the graph connected. That is, F should minimize $\sum_{e \in F} c_e$.

Observation 3

Let $T = (V, F)$ be a minimum-cost solution to the problem described above. Then, T is a tree.

Proof.

- By the definition of the problem, T must be connected.
- By way of contradiction, assume that T has a cycle C . Remove any edge from C resulting in a graph T' . T' is still connect and has a cost less than T .



ALGORITHM DESIGN

TopHat Discussion 3: What greedy heuristic might work?

Kruskal's (1956) Algorithm

- Sort edges by cost from lowest to highest.
- Insert edges unless insertion would create a cycle.

Prim's (1957) Algorithm

- Initialize a node set S with an arbitrary node s .
- Keep the least expensive edge as long as it does not create a cycle.

Reverse-Delete (Kruskal's 1956) Algorithm

- Sort edges by cost from highest to lowest.
- Remove edges unless graph would become disconnected.

ASSUME DISTINCT WEIGHTS

WLOG (WITHOUT LOSS OF GENERALITY)

Theorem 9

(HW Q2) If all edge weights in a connected graph are distinct, then G has a unique MST.

Observation 4

All we need is a consistent tie-breaker when $c_{e_1} = c_{e_2}$ for some pair of edges. I.e. based on the labels of the vertices of $e_1 \cup e_2$.

Assumption: all edge weights are distinct.

ANALYZING MST HEURISTICS

Lemma 10

Let $S \subset V$ be a non-empty proper subset of the nodes, and let $e = (v, w)$ be the minimum cost edge connecting S and $V \setminus S$. Then, every MST contains e .

Proof.

By exchange argument:

- Let T be a spanning tree that does not contain e .
- Let $e' = (v', w')$, where e' is in $P_{v,w} \in T$, $v' \in S$, and $w' \in V \setminus S$.
- Let $T' = T \setminus e' \cup e$.
- T' is connected as e is a $P_{v,w} \in T'$.
- Since $c_e < c_{e'}$, cost of T' is less than T .



KRUSKAL'S ALGORITHM IS OPTIMAL

Kruskal's (1956) Algorithm

- Sort edges by cost from lowest to highest.
- Insert edges unless insertion would create a cycle.

Theorem 11

Kruskal's Algorithm produces an MST.

Proof.

- Let $e = (v, w)$ be the edge added at any step i .
- Since e does not create a cycle, $v \in S$ and $w \notin S$ (WLOG).
- As c_e is the minimum cost edge, the claim follows from Lemma 10.



PRIM'S ALGORITHM IS OPTIMAL

Prim's (1957) Algorithm

- Initialize a node set S with an arbitrary node s .
- Keep the least expensive edge as long as it does not create a cycle.

Theorem 12

Prim's Algorithm produces an MST.

Proof.

- Immediate from Lemma 10.
- That is, Prim's algorithm does exactly what Lemma 10 describes.



REVERSE-DELETE IS OPTIMAL

Reverse-Delete (Kruskal's 1956) Algorithm

- Sort edges by cost from highest to lowest.
- Remove edges unless graph would become disconnected.

How should we prove that it produces an MST?

REVERSE-DELETE IS OPTIMAL

Lemma 13

Let C be any cycle in G , and let e be the most expensive edge of C . Then, e is not in any MST of G .

Proof.

- Let T be a spanning tree that does contain e .
- Let S and $V \setminus S$ be the nodes of the connected components after removing e from T .
- Let e' be an edge in C that connects S and $V \setminus S$.
- Let $T' = T \setminus e \cup e'$.
- T' is connected as e' reconnects S and $V \setminus S$.
- Since $c_e > c_{e'}$, cost of T' is less than T .



REVERSE-DELETE IS OPTIMAL

Lemma 13

Let C be any cycle in G , and let e be the most expensive edge of C . Then, e is not in any MST of G .

Theorem 14

Reverse-Delete Algorithm produces an MST.

Proof.

- Let $e = (v, w)$ be an edge removed at any step i .
- By definition e , belongs to a cycle C .
- As c_e is the maximum cost edge of C , the claim follows from Lemma 13.



IMPLEMENTING PRIM'S ALGORITHM

Prim's (1957) Algorithm

- Initialize a node set S with an arbitrary node s .
- Keep the least expensive edge as long as it does not create a cycle.

Key Operations

- Retrieve the minimum valued edge between S and $V \setminus S$.
- Prim's and Dijkstra's have nearly identical implementations (but different minimizers)!

Priority Queue (min-heap)

- ExtractMin ($O(1)$): $n - 1$ times.
- ChangeKey ($O(\log(n))$): m times.

Overall: $O(m \log(n))$

IMPLEMENTING KRUSKAL'S ALGORITHM

Kruskal's (1956) Algorithm

- Sort edges by cost from lowest to highest.
- Insert edges unless insertion would create a cycle.

Key Operations

- Sorting the edges: ($O(m \log m)$ and, since $m \leq n^2$, $O(m \log n)$).
- Maintain sets of connected components that we merge.
- Initialize one set per node: $O(n)$.

Union-Find Data Structure

- Find(x): Finds the set containing x . ($O(\log n)$ can be $O(\alpha(n))$)
- Union(x, y): Joins two sets x and y . ($O(1)$)

UNION-FIND / DISJOINT-SET

Key Operations


- Find(x): Finds the set containing x . ($O(\log n)$ can be $O(\alpha(n))$)
- Union(x, y): Joins two sets x and y . ($O(1)$)

Basic Container

node	rank	parent
------	------	--------

Initializing Data Structure for Kruskal's

For each node s , create a singleton set. That is each container has rank 0 and points to itself.



UNION-FIND OPERATIONS

Find(x): $O(\log n)$

- If $x.\text{parent}$ points to x , return x .
- Else Find($x.\text{parent}$)
- $O(\log n)$ requires balanced trees.
- $O(\alpha(n))$ with *path compression*.

Union(x, y): $O(1)$

- (WLOG) $x.\text{rank} \geq y.\text{rank}$:
 $y.\text{parent} = x$
- If $x.\text{rank} = y.\text{rank}$:
 $x.\text{rank} := x.\text{rank} + 1$
- By using rank, we maintain balanced sets if we start with balanced sets.

IMPLEMENTING KRUSKAL'S ALGORITHM

Kruskal's (1956) Algorithm

- Sort edges by cost from lowest to highest.
- Insert edges unless insertion would create a cycle.

Key Operations

- Sorting the edges: ($O(m \log m)$ and, since $m \leq n^2$, $O(m \log n)$).
- Maintain sets of connected components that we merge.
- Initialize one set per node: $O(n)$.

Union-Find Data Structure

TH: How many Find and Unions?

- Find(x): Finds the set containing x .
- Union(x, y): Joins two sets x and y .

IMPLEMENTING KRUSKAL'S ALGORITHM

Kruskal's (1956) Algorithm

- Sort edges by cost from lowest to highest.
- Insert edges unless insertion would create a cycle.

Key Operations

- Sorting the edges: ($O(m \log m)$ and, since $m \leq n^2$, $O(m \log n)$).
- Maintain sets of connected components that we merge.
- Initialize one set per node: $O(n)$.

Union-Find Data Structure

- Find(x): $2m$ times $O(\log n)$ (can be $O(\alpha(n))$).
- Union(x,y): $n - 1$ times $O(1)$.

GRAPH EXPLORATION OVERVIEW

BFS and DFS

- Traverses a graph G starting from some node s .
- Builds a tree T .
- No guarantee on any distance measure.

Dijkstra's


- Traverses a graph starting from some node s .
- Builds a tree T .
- All s to u paths in T are the shortest such path in G .

MST Algorithms

- Explores a graph G edges.
- Builds a tree T .
- T is minimum cost to connect all nodes in G .

CLUSTERING

k -CLUSTERING



Maximizing Spacing Problem

- A universe $\mathcal{U} := \{p_1, \dots, p_n\}$ of n objects.
- Distance function $d : \mathcal{U} \times \mathcal{U} \rightarrow \mathbb{R}$ such that, for all $p_i, p_j \in \mathcal{U}$:
 - $d(p_i, p_i) = 0$
 - $d(p_i, p_j) > 0$
 - $d(p_i, p_j) = d(p_j, p_i)$
- Objective: Partition \mathcal{U} into k non-empty groups $\mathcal{C} := C_1, \dots, C_k$ with maximum spacing:

$$\text{maximize } \min_{C_i, C_j \in \mathcal{C}} \min_{u \in C_i, v \in C_j} d(u, v)$$

ALGORITHM DESIGN

TopHat Discussion 4: What greedy approach might work?

ALGORITHM DESIGN

Algorithm

- Build an MST.
- Remove $k - 1$ largest edges.

k -Clusters at max spacing?

- Start with a tree, remove $k - 1$ edges: We get a forest of k trees.
- By definition largest edges are removed so max spacing.

TopHat Q10: Which MST algorithm?

Kruskal's ($O(m \log n)$ which is $O(n^2 \log n)$ for clustering):

- Merge sets from lowest to most expensive edges.
- Stop when we have k sets.

PREFIX CODES

BINARY ENCODING

Fixed-Width Encoding

- Set of symbols $S := \{a, b, c, d, e\}$.
- Encoding function $\gamma : S \rightarrow \{0, 1\}^k$.
 $\gamma(S) := \{000, 001, 010, 011, 100\}$.
- Ex. ASCII
- TopHat Q11: Decode 000010.

Variable-Width Encoding

- Set of symbols $S := \{a, b, c, d, e\}$.
- Encoding function $\gamma : S \rightarrow \{0, 1\}^*$.
 $\gamma(S) := \{0, 1, 10, 01, 11\}$.
- TopHat Q12: How many ways to decode 0010?

UNIQUE VARIABLE-WIDTH ENCODINGS

Prefix Codes

Encoding of S such that no encoding of a symbol in S is a prefix of another.

- Set of symbols $S := \{a, b, c, d, e\}$.
- Encoding function $\gamma : S \rightarrow \{0, 1\}^*$.
 $\gamma(S) := \{11, 01, 001, 000, 100\}$.
- 0010 invalid sequence
- TopHat 13: Decode 1101.

UNIQUE VARIABLE-WIDTH ENCODINGS

Prefix Codes

Encoding of S such that no encoding of a symbol in S is a prefix of another.

- Set of symbols $S := \{a, b, c, d, e\}$.
- Encoding function $\gamma : S \rightarrow \{0, 1\}^*$.
 $\gamma(S) := \{11, 01, 001, 000, 100\}$.

Easy Decoding



Scan left to right, once an encoding is matched, output symbol.

Optimal Prefix Codes

- For a set of symbols S , let f_x denote the frequency of x in the text to be encoded.
- Average bits $\text{ABL}(\gamma) := \sum_{x \in S} f_x \cdot |\gamma(x)|$.
- Goal: Find γ that minimizes ABL .

ALGORITHM DESIGN

PREFIX BINARY TREES



OPTIMAL PREFIX TREE IS FULL

Theorem 15

The binary tree corresponding to the optimal prefix code is full.

Proof.

By exchange argument:

- Let T be an optimal prefix tree with a node u with one child v .
- Let T' be T with u replaced with v .
- Distance to v decreases by 1 in T' , a contradiction.



TOP-DOWN APPROACH

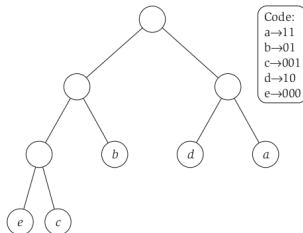
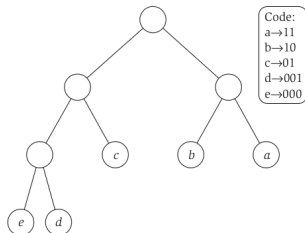
Algorithm

- Split S into two sets such that the sets frequency are $1/2$ the total frequency.
- Recurse on new sets until singletons.

$$f_a = .32, f_b = .25, f_c = .2, f_d = .18, f_e = .05$$

$$\text{ABL}(\text{OPT}) = 2.23$$

$$\text{ABL}(\text{TopDown}) = 2.25$$



WHAT IF WE KNEW THE OPTIMAL TREE?

Let T^* be the optimal (unlabelled) prefix tree.

Lemma 16

Let u, v be leaves of T^ such that $\text{depth}(u) < \text{depth}(v)$, where u is labelled with y and v is labelled with z . Then, $f_y \geq f_z$.*

Proof.

If $f_y < f_z$, exchange the labelling of y and z . Since $\text{depth}(u) < \text{depth}(v)$, $\text{ABL}(T^*)$ must decrease with the new labelling. □

WHAT IF WE KNEW THE OPTIMAL TREE?

Let T^* be the optimal (unlabelled) prefix tree.

Lemma 16

Let u, v be leaves of T^ such that $\text{depth}(u) < \text{depth}(v)$, where u is labelled with y and v is labelled with z . Then, $f_y \geq f_z$.*

Labelling T^*

- Order symbols by increasing frequency.
- Assign them to leaves of T^* by decreasing depth.

Observation 5


In T^ , the lowest frequency letters are siblings.*

BOTTOM-UP APPROACH

HUFFMAN CODE

Huffman's Algorithm

- (1) Bottom-up by lowest frequency:
 - Let x and y be the lowest frequency symbols.
 - Set $S := S \setminus \{x, y\} \cup \{w := xy\}$ and $f_w = f_x + f_y$.
 - Repeat until $|S| = 1$.
- (2) Generate the tree:
 - $T :=$ root with element from S .



- Replace
- Repeat until leaves of T are original symbols.

HUFFMAN CODES ARE OPTIMAL

Lemma 17

Let T' be the tree at the $(k - 1)$ -st step, and let T be the tree at the k -th step. $ABL(T') = ABL(T) - f_w$, where w is the symbol replaced in the k -th step by y and z .

Proof.

$$\begin{aligned} ABL(T) &= \sum_{x \in S} f_x \cdot \text{depth}(x) \\ &= f_y \cdot \text{depth}(y) + f_z \cdot \text{depth}(z) + \sum_{x \in S; x \notin \{y, z\}} f_x \cdot \text{depth}(x) \\ &= f_w + f_w \cdot \text{depth}(w) + \sum_{x \in S \setminus \{y, z\}} f_x \cdot \text{depth}(x) \\ &= f_w + ABL(T') \end{aligned}$$



HUFFMAN CODES ARE OPTIMAL

Lemma 17

Let T' be the tree at the $(k - 1)$ -st step, and let T be the tree at the k -th step. $ABL(T') = ABL(T) - f_w$, where w is the symbol replaced in the k -th step by y and z .

Theorem 18

Huffman Algorithm is optimal.

Proof.

By induction:

- Base case $|S| = 2$
- Inductive step: We have T . By way of contradiction, assume $ABL(Z) \leq ABL(T)$.

HUFFMAN CODES ARE OPTIMAL

Lemma 17

Let T' be the tree at the $(k - 1)$ -st step, and let T be the tree at the k -th step. $ABL(T') = ABL(T) - f_w$, where w is the symbol replaced in the k -th step by y and z .

Theorem 18

Huffman Algorithm is optimal.

Proof.

By induction:

- We observed that y and z are siblings. Hence:

$$ABL(Z) < ABL(T)$$

$$\iff ABL(Z') + f_w < ABL(T') + f_w, \text{ by Lemma 17}$$


$$\iff ABL(Z') < ABL(T'), \text{ a contradiction.} \quad \square$$

BOTTOM-UP APPROACH

HUFFMAN CODE

Huffman's Algorithm

- (1) Bottom-up by lowest frequency:
 - Let x and y be the lowest frequency symbols.
 - Set $S := S \setminus \{x, y\} \cup \{w := xy\}$ and $f_w = f_x + f_y$.
 - Repeat until $|S| = 1$.
- (2) Generate the tree:
 - $T :=$ root with element from S .



- Replace
- Repeat until leaves of T are original symbols.


Runtime: $|S| - 1$ recursions with find min over $|S_i|$ elements

BOTTOM-UP APPROACH

HUFFMAN CODE

Huffman's Algorithm

- (1) Bottom-up by lowest frequency:
 - Let x and y be the lowest frequency symbols.
 - Set $S := S \setminus \{x, y\} \cup \{w := xy\}$ and $f_w = f_x + f_y$.
 - Repeat until $|S| = 1$.
- (2) Generate the tree:
 - $T :=$ root with element from S .



- Replace
- Repeat until leaves of T are original symbols.

Runtime: $O(|S|^2)$

what about $O(|S| \log |S|)$? Priority Queue (min-heap)