

Answer the questions in the boxes provided on the question sheets. If you run out of room for an answer, add a page to the end of the document.

Name: \_\_\_\_\_ Wisc id: \_\_\_\_\_

## Greedy Algorithms

1. In one or two sentences, describe what a greedy algorithm is. Your definition should be informal, something you could share with a non computer scientist.

2. There are many different problems all described as “scheduling” problems. In the following questions, pay attention to the details of the problem setup, as they will change each time!

- (a) Let each job have a start time, an end time, and a value. We want to schedule as much value of non-conflicting jobs as possible. Use a counterexample to show that Earliest Finish First (the greedy algorithm we used for jobs with all equal value) does NOT work in this case.

- (b) *Kleinberg, Jon. Algorithm Design (p. 191, q. 7)* Now let each job consist of two durations. A job  $i$  must be preprocessed for  $p_i$  time on a supercomputer, and then finished for  $f_i$  time on a standard PC. There are enough PCs available to run all jobs at the same time, but there is only one supercomputer (which can only run a single job at a time). The completion time of a schedule is defined as the earliest time when all jobs are done running on both the supercomputer and the PCs. Give a polynomial time algorithm that finds a schedule with the earliest completion time possible.

- (c) Prove the correctness and efficiency of your algorithm from part (b).

3. Kleinberg, Jon. *Algorithm Design* (p. 190, q. 5)

- (a) Consider a long straight road with houses scattered along it. We want to place cell phone towers along the road so that every house is within four miles of at least one tower. Give an efficient algorithm that achieves this goal using the minimum possible number of towers.

- (b) Prove the correctness of your algorithm.

4. *Kleinberg, Jon. Algorithm Design (p. 197, q. 18)* Your friends are planning to drive north from Madison to the town of Superior, Wisconsin over winter break. They have drawn a directed graph with nodes representing potential stops and edges representing the roads between them.

They have also found a weather forecasting site that can accurately predict how long it will take to traverse one of the edges on their graph, given the starting time  $t$ . This is important because some of the roads on their graph are affected strongly by the seasons and by extreme weather. It's guaranteed that it never takes negative time to traverse an edge, and that you can never arrive earlier by starting later.

- (a) Design an algorithm your friends can use to plot the quickest route. You may assume that they start at time  $t = 0$ , and that the predictions made by the weather forecasting site are accurate.

- (b) Demonstrate how your algorithm works using a small example with 6 nodes. Your demonstration should include any data structures you maintain during the execution of your algorithm and any queries you make to the weather forecasting site. For example, if your algorithm maintains a “current path” that grows from (M)adison to (S)uperior, you might show something like the following table:

| Path      | Total time |
|-----------|------------|
| M         | 0          |
| M,A       | 2          |
| M,A,E     | 5          |
| M,A,E,F   | 6          |
| M,A,E     | 5          |
| M,A,E,H   | 10         |
| M,A,E,H,S | 13         |

5. *Kleinberg, Jon. Algorithm Design (p. 189, q. 3).*

You are consulting for a trucking company that does a large amount of business shipping packages between New York and Boston. The volume is high enough that they have to send a number of trucks each day between the two locations. Trucks have a fixed limit  $W$  on the maximum amount of weight they are allowed to carry. Boxes arrive at the New York station one by one, and each package  $i$  has a weight  $w_i$ . The trucking station is quite small, so at most one truck can be at the station at any time. Company policy requires that boxes are shipped in the order they arrive; otherwise, a customer might get upset upon seeing a box that arrived after his make it to Boston faster. At the moment, the company is using a simple greedy algorithm for packing: they pack boxes in the order they arrive, and whenever the next box does not fit, they send the truck on its way.

Prove that, for a given set of boxes with specified weights, the greedy algorithm currently in use actually minimizes the number of trucks that are needed. Hint: Use the stay ahead method.

6. *Kleinberg, Jon. Algorithm Design (p. 192, q. 8).* Suppose you are given a connected graph  $G$  with edge costs that are all distinct. Prove that  $G$  has a unique minimum spanning tree.

7. *Kleinberg, Jon. Algorithm Design (p. 193, q. 10).* Let  $G = (V, E)$  be an (undirected) graph with costs  $c_e \geq 0$  on the edges  $e \in E$ . Assume you are given a minimum-cost spanning tree  $T$  in  $G$ . Now assume that a new edge is added to  $G$ , connecting two nodes  $v, w \in V$  with cost  $c$ .
- (a) Give an efficient ( $O(|E|)$ ) algorithm to test if  $T$  remains the minimum-cost spanning tree with the new edge added to  $G$  (but not to the tree  $T$ ). Please note any assumptions you make about what data structure is used to represent the tree  $T$  and the graph  $G$ , and prove that its runtime is  $O(|E|)$ .

- (b) Suppose  $T$  is no longer the minimum-cost spanning tree. Give a linear-time algorithm (time  $O(|E|)$ ) to update the tree  $T$  to the new minimum-cost spanning tree. Prove that its runtime is  $O(|E|)$ .

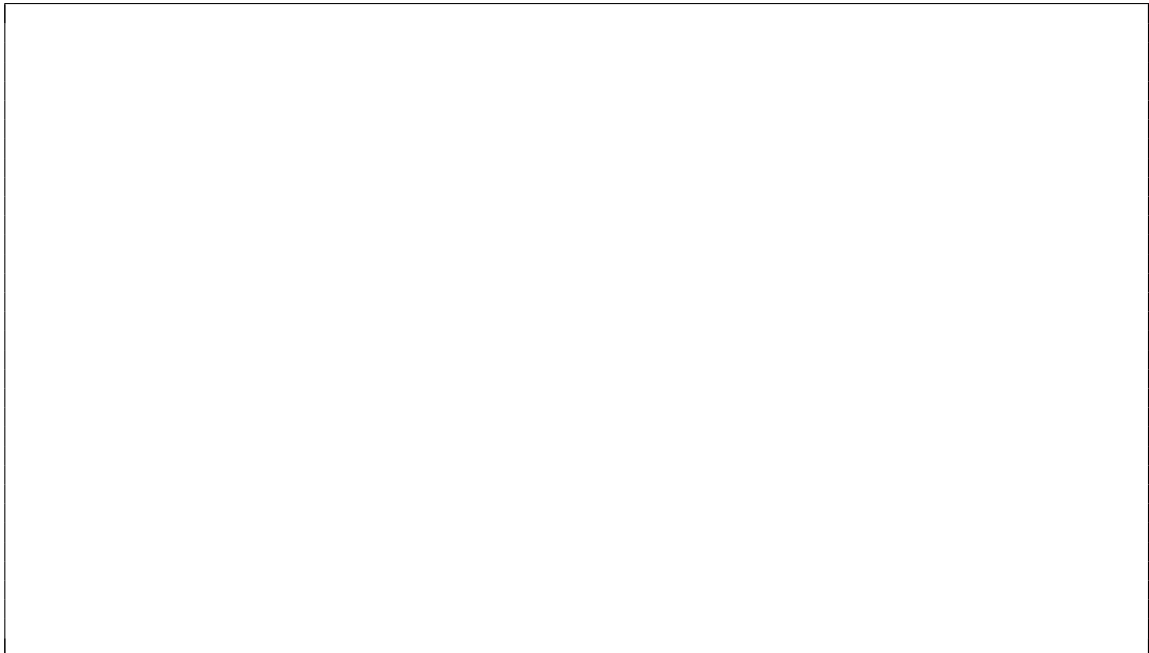


8. In class, we saw that an optimal greedy strategy for the paging problem was to reject the page the furthest in the future (FF). The paging problem is a classic online problem, meaning that algorithms do not have access to future requests. Consider the following online eviction strategies for the paging problem, and provide counter-examples that show that they are not optimal offline strategies.<sup>1</sup>

(a) FWF is a strategy that, on a page fault, if the cache is full, it evicts all the pages.



(b) LRU is a strategy that, if the cache is full, evicts the least recently used page when there is a page fault.



---

<sup>1</sup>An interesting note is that both of these strategies are  $k$ -competitive, meaning that they are equivalent under the standard theoretical measure of online algorithms. However, FWF really makes no sense in practice, whereas LRU is used in practice.

## Coding Questions

### 9. Interval Scheduling:

Implement the optimal algorithm for interval scheduling (for a definition of the problem, see the Greedy slides on Canvas) in either C, C++, C#, Java, Python, or Rust. Be efficient and implement it in  $O(n \log n)$  time, where  $n$  is the number of jobs.

The input will start with a positive integer, giving the number of instances that follow. For each instance, there will be a positive integer, giving the number of jobs. For each job, there will be a pair of positive integers  $i$  and  $j$ , where  $i < j$ , and  $i$  is the start time, and  $j$  is the end time.

A sample input is the following:

```
2
1
1 4
3
1 2
3 4
2 6
```

The sample input has two instances. The first instance has one job to schedule with a start time of 1 and an end time of 4. The second instance has 3 jobs.

For each instance, your program should output the number of intervals scheduled on a separate line. Each output line should be terminated by a newline. The correct output to the sample input would be:

```
1
2
```

### 10. Paging:

For this question you will implement Furthest in the future paging in either C, C++, C#, Java, Python, or Rust.

The input will start with a positive integer, giving the number of instances that follow. For each instance, the first line will be a positive integer, giving the number of pages in the cache. The second line of the instance will be a positive integer giving the number of page requests. The third and final line of each instance will be space delimited positive integers which will be the request sequence.

Note: a naïve solution doing repeated linear searches will timeout.

A sample input is the following:

```
3
2
7
1 2 3 2 3 1 2
4
12
12 3 33 14 12 20 12 3 14 33 12 20
3
20
1 2 3 4 5 1 2 3 4 5 1 2 3 4 5 1 2 3 4 5
```

The sample input has three instances. The first has a cache which holds 2 pages. It then has a request sequence of 7 pages. The second has a cache which holds 4 pages and a request sequence of 12 pages. The third has a cache which holds 3 pages and a request sequence of 15 pages.

For each instance, your program should output the number of page faults achieved by furthest in the future paging assuming the cache is initially empty at the start of processing the page request sequence. One output should be given per line. The correct output for the sample input is

4  
6  
12