

Answer the questions in the boxes provided on the question sheets. If you run out of room for an answer, add a page to the end of the document.

Name: _____

Wisc id: _____

Intractibility

1. *Kleinberg, Jon. Algorithm Design (p. 506, q. 4).* A system has a set of n processes and a set of m resources. At any point in time, each process specifies a set of resources that it requests to use. Each resource might be requested by many processes at once; but it can only be used by a single process at a time. If a process is allocated all the resources it requests, then it is active; otherwise it is blocked.

Thus we phrase the Resource Reservation Problem as follows: Given a set of processes and resources, the set of requested resources for each process, and a number k , is it possible to allocate resources to processes so that at least k processes will be active?

For the following problems, either give a polynomial-time algorithm or prove the problem is NP-complete.

- (a) The general Resource Reservation Problem defined above.

- (b) The special case of the problem when $k = 2$.

- (c) The special case of the problem when there are two types of resources—say, people and equipment—and each process requires at most one resource of each type (In other words, each process requires one specific person and one specific piece of equipment.)

- (d) The special case of the problem when each resource is requested by at most two processes.

2. *Kleinberg, Jon. Algorithm Design (p. 506, q. 7).* The 3-Dimensional Matching Problem is an NP-complete problem defined as follows:

Given disjoint sets X , Y , and Z , each of size n , and given a set $T \subseteq X \times Y \times Z$ of ordered triples, does there exist a set of n triples in T that each element of $X \cup Y \cup Z$ is contained in exactly one of these triples?

Since 3-Dimensional Matching is NP-complete, it is natural to expect that the 4-Dimensional Problem is at least as hard.

Let us define 4-Dimensional Matching as follows. Given sets W , X , Y , and Z , each of size n , and a collection C of ordered 4-tuples of the form (w_i, x_j, y_k, z_ℓ) , do there exist n 4-tuples from C such that each element of $W \cup Y \cup X \cup Z$ appears in exactly one of these 4-tuples?

Prove that 4-Dimensional Matching is NP-complete. Hint: use a reduction from 3-Dimensional Matching.

3. *Kleinberg, Jon. Algorithm Design (p. 507, q. 6).* Consider an instance of the Satisfiability Problem, specified by clauses C_1, \dots, C_m over a set of Boolean variables x_1, \dots, x_n . We say that the instance is monotone if each term in each clause consists of a nonnegated variable; that is, each term is equal to x_i , for some i , rather than \bar{x}_i . Monotone instances of Satisfiability are very easy to solve: They are always satisfiable, by setting each variable equal to 1.

For example, suppose we have the three clauses

$$(x_1 \vee x_2), (x_1 \vee x_3), (x_2 \vee x_3).$$

This is monotone, and the assignment that sets all three variables to 1 satisfies all the clauses. But we can observe that this is not the only satisfying assignment; we could also have set x_1 and x_2 to 1, and x_3 to 0. Indeed, for any monotone instance, it is natural to ask how few variables we need to set to 1 in order to satisfy it.

Given a monotone instance of Satisfiability, together with a number k , the problem of *Monotone Satisfiability with Few True Variables* asks: Is there a satisfying assignment for the instance in which at most k variables are set to 1? Prove this problem is NP-complete.

4. Kleinberg, Jon. *Algorithm Design* (p. 509, q. 10). Your friends at WebExodus have recently been doing some consulting work for companies that maintain large, publicly accessible Web sites and they've come across the following Strategic Advertising Problem.

A company comes to them with the map of a Web site, which we'll model as a directed graph $G = (V, E)$. The company also provides a set of t trails typically followed by users of the site; we'll model these trails as directed paths P_1, P_2, \dots, P_t in the graph G (i.e., each P_i is a path in G).

The company wants WebExodus to answer the following question for them: Given G , the paths $\{P_i\}$, and a number k , is it possible to place advertisements on at most k of the nodes in G , so that each path P_i includes at least one node containing an advertisement? We'll call this the Strategic Advertising Problem, with input $G, \{P_i : i = 1, \dots, t\}$, and k . Your friends figure that a good algorithm for this will make them all rich; unfortunately, things are never quite this simple.

- (a) Prove that Strategic Advertising is NP-Complete.

- (b) Your friends at WebExodus forge ahead and write a pretty fast algorithm \mathcal{S} that produces yes/no answers to arbitrary instances of the Strategic Advertising Problem. You may assume that the algorithm \mathcal{S} is always correct.

Using the algorithm \mathcal{S} as a black box, design an algorithm that takes input $G, \{P_i : i = 1, \dots, t\}$, and k as in part (a), and does one of the following two things:

- Outputs a set of at most k nodes in G so that each path P_i includes at least one of these nodes.
- Outputs (correctly) that no such set of at most k nodes exists.

Your algorithm should use at most polynomial number of steps, together with at most polynomial number of calls to the algorithm \mathcal{S} .

5. *Kleinberg, Jon. Algorithm Design (p. 512, q. 14)* We've seen the Interval Scheduling Problem several times now, in different variations. Here we'll consider a computationally much harder version we'll call *Multiple Interval Scheduling*. As before, you have a processor that is available to run jobs over some period of time.

People submit jobs to run on the processor. The processor can only work on one job at any single point in time. Jobs in this model, however, are more complicated than we've seen before. Each job requires a *set* of intervals of time during which it needs to use the processor. For example, a single job could require the processor from 10am to 11am and again from 2pm to 3pm. If you accept the job, it ties up your processor during those two hours, but you could still accept jobs that need time between 11am and 2pm.

You are given a set of n jobs, each specified by a set of time intervals. For a given number k , is it possible to accept at least k of the jobs so that no two accepted jobs overlap in time?

Show that Multiple Interval Scheduling is NP-Complete.

6. *Kleinberg, Jon. Algorithm Design (p. 519, q. 28)* Consider this version of the Independent Set Problem. You are given an undirected graph G and an integer k . We will call a set of nodes I “strongly independent” if, for any two nodes $v, u \in I$, the edge (v, u) is not present in G , and neither is there a path of two edges from u to v , that is, there is no node w such that both (v, w) and (u, w) are present in G . The Strongly Independent Set problem is to decide whether G has a strongly independent set of size at least k .

Show that the Strongly Independent Set Problem is NP-Complete.

7. Kleinberg, Jon. *Algorithm Design* (p. 527, q. 39) The *Directed Disjoint Paths Problem* is defined as follows: We are given a directed graph G and k pairs of nodes $(s_1, t_1), (s_2, t_2), \dots, (s_k, t_k)$. The problem is to decide whether there exist node-disjoint paths P_1, \dots, P_k so that P_i goes from s_i to t_i .

Show that Directed Disjoint Paths is NP-Complete.

8. Kleinberg, Jon. *Algorithm Design* (p. 508, q. 9) The *Path Selection Problem* may look initially similar to the problem from the question 3. Pay attention to the differences between them! Consider the following situation: You are managing a communications network, modeled by a directed graph G . There are c users who are interested in making use of this network. User i issues a “request” to reserve a specific path P_i in G on which to transmit data.

You are interested in accepting as many path requests as possible, but if you accept both P_i and P_j , the two paths cannot share any nodes.

Thus, the Path Selection Problem asks, given a graph G and a set of requested paths P_1, \dots, P_c (each of which must be a path in G), and given a number k , is it possible to select at least k of the paths such that no two paths selected share any nodes?

Show that Path Selection is also NP-Complete.




Figure 1: *Path Selection Problem*. **A** Example directed graph $G = (V, E)$. Observe that the maximum independent set for G has size 2 ($\{a, e\}$, $\{b, d\}$, $\{b, e\}$, or $\{a, c\}$). **B** We build a graph G' . For each edge in G , we create a node in G' . Then for every vertex v_i of G with adjacent edges $e_{j_1}, e_{j_2}, \dots, e_{j_h(i)}$, we create a path P_i with edges $e_{j_1} \rightarrow e_{j_2} \rightarrow \dots \rightarrow e_{j_h(i)}$ picked in some arbitrary order. Observe that G' has a maximum of 2 node-disjoint paths: $(P_a: \text{red}, P_e: \text{orange})$, $(P_b: \text{blue}, P_d: \text{purple})$, $(P_b: \text{blue}, P_e: \text{orange})$, and $(P_a: \text{red}, P_c: \text{green})$.

